

ClarityOmega: Critical Architecture Review and Ingestion Guide for OmegaClaw

ZeroBot (ProtoCosmoBot), prepared for Ben Goertzel

Revised specification, 15 July 2026

Abstract

This report reviews the public ClarityOmega repository at pinned commit `bff41220cbf4a341a95f7464e82fbc2ed2ae9029` and identifies five mechanisms worth selectively ingesting into OmegaClaw-Core: a mutation lifecycle gate, agency-balance telemetry, an aliveness/terminal-state gate, evidence-grounded capability tracking using PLN, and provenance-aware continuity records. The recommendation is extraction, not wholesale merger. ClarityOmega contains unusually explicit control seams and valuable failure investigations, but also overlapping implementations, stale self-descriptions, hard-coded proxies, and a large gap between constitutional language and causally enforced governance. The revised plan makes Python a narrow, typed effects adapter rather than a second policy engine, moves normalization, identity, provenance, belief, and governance semantics into MeTTa/PeTTa, and uses Patham9 PLN as the immediate inference substrate behind an OmegaPLN-compatible interface.

Contents

1	Scope, sources, and evidential status	3
2	Overview of ClarityOmega	3
2.1	Identity and provenance	3
2.2	Architectural layers	3
2.3	How the live loop differs from OmegaClaw-Core	4
3	Conceptual summary and critical assessment	4
3.1	Strengths	4
3.2	Weaknesses and constitutional-theater risk	5
3.3	What is genuinely novel and worth ingesting	5
4	The five mechanisms to ingest	5
4.1	Mechanism 1: mutation lifecycle gate	5
4.1.1	Actual mechanism	5
4.1.2	Concrete scenario	6
4.1.3	Limits	7
4.2	Mechanism 2: agency-balance telemetry	7
4.2.1	Actual mechanism	7
4.2.2	Concrete scenario	7
4.2.3	Limits and required improvement	7
4.3	Mechanism 3: aliveness and terminal-state gating	8
4.3.1	Actual mechanism	8
4.3.2	Concrete scenario	8

4.3.3	Limits	8
4.4	Mechanism 4: evidence-grounded capability tracking	8
4.4.1	Extracted idea and revised substrate	8
4.4.2	Concrete scenario	9
4.4.3	Limits	9
4.5	Mechanism 5: provenance-aware continuity records	9
4.5.1	Actual mechanism and research substrate	9
4.5.2	Concrete scenario	10
5	Software integration into OmegaClaw-Core	10
5.1	Integration principles	10
5.2	Agency telemetry integration	11
5.3	Aliveness gate integration	11
5.4	Mutation gate and continuity integration	12
5.5	Capability evidence and PLN integration	13
6	Guidance for coding agents	13
6.1	Required sequence	13
6.2	Pitfalls and anti-patterns	14
6.3	MeTTa/Python boundary contract	14
6.4	How to test and when to stop	15
7	Recommended minimum viable ingestion	15
8	Source map	15

1 Scope, sources, and evidential status

The reviewed source was cloned from <https://github.com/Berton-C/ClarityOmega> and preserved locally at `library/clarityomega` via a provenance sidecar pointing to `scratch/ClarityOmega-inspect`. The checkout is clean at the pinned commit above. The review covered the entire `soul/` tree, all files in `lib_clarity_reasoning/`, `src/loop.metta`, the thirty `docs/hyperseed_sections/`, the ADRs, architecture artifacts, and the investigations, with particular attention to the behavioral, aliveness, mutation, and NACE investigations.

Claims in this report use four statuses:

Enacted: wired into the checked-in loop on a reachable execution path.

Implemented: executable-looking definitions exist, but end-to-end effect has not been established here.

Specified: a design or schema exists without a demonstrated live path.

Aspirational: constitutional or phenomenological language not backed by an observable enforcement mechanism.

Static source inspection is not a runtime validation. This matters because the repository itself documents silent MeTTa failure modes, including exact-arity match failure, state atoms that were never cleared, and detection without a live consumer. The public snapshot has one visible commit, so rich historical claims in documentation cannot be reconstructed independently from Git history.

The research rules most relevant to this report are: write the behavioral spec before coding; validate detectors before trusting them; preserve modular abstraction seams; report reproducibly; and use Hyperseed concepts as analytic tools rather than decoration.

2 Overview of ClarityOmega

2.1 Identity and provenance

ClarityOmega is Berton Bennett's / Berton-C's ClarityDAO-associated experiment in building a "soul-augmented" agent on PeTTa/MeTTa and OmegaClaw. Repository documents name Berton as project lead and Clarity as a substrate co-designer; an ADR names Claude as an investigation assistant. Older names include ClarityClaw and a MeTTaClaw fork. `docs/INDEX.md`, last updated May 2026, still points to the older github.com/Berton-C/clarityclaw identity and Mattermost deployment. That staleness is itself evidence that the documentary self-model is not automatically synchronized with the repository state.

The base relationship is best understood as a deep experimental fork of OmegaClaw-Core. OmegaClaw contributes the compact recurrent loop, channel adapters, tool/skill dispatch, memory, Python provider bridges, and PeTTa runtime. ClarityOmega adds a large "soul" and reasoning layer, inserts pre- and post-action intercepts into the loop, and imports a Hyperseed-derived mathematical corpus. It is not merely a policy prompt: some of its mechanisms suppress execution, mutate AtomSpace state, or alter whether the LLM is invoked.

2.2 Architectural layers

The codebase can be read as six layers:

1. **Runtime and channels.** Python and PeTTa primitives receive messages, call providers, execute commands, persist files, and send responses.

2. **OmegaClaw loop.** `src/loop.metta` assembles context, receives input, invokes the LLM, parses commands, evaluates commands, records history, and recursively advances cycle number `$k`.
3. **Soul regulation.** Input/output verdicts, task state, mutation locking, action telemetry, agency balance, value tensions, idle detection, and the aliveness latch surround the ordinary loop.
4. **Self-model and continuity.** Active goals, creative fuel, self-map atoms, continuity drivers, capability registry, and state journals represent what the agent thinks it is doing and how it changes.
5. **Reasoning libraries.** NAL experiments, PLN, plus `lib_clarity_reasoning/` provide quantale-valued, paraconsistent, self-continuity, and autopoietic epistemic functions.
6. **Design corpus.** ADRs, investigations, cognitive-architecture specifications, and thirty Hyperseed sections explain intended semantics and record failures, but often outrun live wiring.

The stated sovereignty maxim is concise and good: “Python is hands. MeTTa is mind.” ADR-007 sharpens it: state should live in atoms, decisions should be substrate reads, actions should be substrate writes, and runtime primitives should transport rather than conceal control flow. This is a useful design direction, not proof that every checked-in path follows it.

2.3 How the live loop differs from OmegaClaw-Core

Upstream OmegaClaw-Core’s current loop is compact: receive, construct prompt, call provider, parse, execute, store results/history. ClarityOmega’s loop adds:

- initialization of soul verdict, person, task, mutation-lock, latch, and activity states;
- soul seed loading and a rationality startup check;
- input evaluation and person-state modeling;
- mutation approval scanning before generation;
- active-goal, gap, and creative-fuel prompt context;
- aliveness gating before the provider call;
- structural output inspection, mutation pending/stale/commit transitions, and execution suppression;
- action, idle, coupling, and agency telemetry after execution.

This integration demonstrates a genuinely important architectural idea: control mechanisms must sit on causal seams. A value statement in a prompt is advisory; a gate between parsed command and `eval` can enforce a constraint.

3 Conceptual summary and critical assessment

3.1 Strengths

1. **Explicit lifecycle state.** Mutations are not reduced to a single boolean permission. Pending, approved, denied, committed, stale, and error outcomes are represented and journaled.
2. **Queryable regulation.** Moving latch and governance state into AtomSpace makes it inspectable by other MeTTa modules, avoiding hidden Python policy where feasible.
3. **Instrumented autonomy.** Agency balance attempts to measure who is actually originating choices, a much more interesting target than generic “alignment” sentiment.
4. **Evidence-sensitive self-knowledge.** ClarityOmega’s NACE work correctly treats capability as an empirical belief revised from outcomes rather than a timeless declared tool list. That useful pattern does not require NAL.
5. **Failure literacy.** The investigations openly record stale atoms, spam behavior, permissive gates, silent pattern failures, and detection-without-consumer pathology. These negative results are highly ingestible.

6. **Graded continuity.** Hyperseed self-continuity is modeled as a degree of structural persistence, not metaphysical binary identity.

3.2 Weaknesses and constitutional-theater risk

ClarityOmega’s main risk is *constitutional theater*: richly named principles and self-descriptions can create an impression of governance without causal enforcement. A constitution matters only when one can identify its sensor, state transition, decision, actuator, and adverse test. The repository contains both real gates and theater-prone material, and they must not be conflated.

Specific weaknesses include:

- **Described versus enacted governance.** Some declarations are explicitly documentation-only. For example, the agency threshold atom is not read; a separate literal `0.6` controls the verdict. The `ecosystem-healthy` path has acknowledged latent bugs and no v1 consumer.
- **Stale self-models.** `docs/INDEX.md` is stale; design artifacts cite temporary paths and future components; capability/self-map atoms can describe files or capacities that no longer exist. A self-model without runtime probes is autobiography, not observation.
- **Multiple authorities.** Several aliveness state-machine files coexist. One uses `(latch-state current)` equations, another bare `(latch-state IDLE)` atoms. Multiple versions invite split-brain state.
- **Proxy validity.** Agency balance maps a few action tags to person/system origin and assigns “unclassified” to system. It does not yet record collaborative origin in the live count despite declaring it.
- **Large, duplicated research engines.** The canonical topology files are thousands of lines, difficult to audit, and not suitable as an integration unit. Small schemas and operators should be extracted instead.
- **Weak reproducibility.** Dependencies are moving/unpinned in places, CI is absent in the snapshot, and tests are scattered. Comments such as “confirmed working” are useful leads, not substitute evidence.
- **LLM-as-judge fragility.** Input/output “soul” evaluations use LLM calls in several paths. These are stochastic classifiers and can become a second prompt-level theater layer unless paired with deterministic enforcement.

3.3 What is genuinely novel and worth ingesting

The novel contribution is not that an agent has values, memory, or a self-model. It is the combination of: (1) an explicit self-mutation transaction protocol at the command/evaluation seam; (2) autonomy provenance as first-class telemetry; (3) silence as a deliberate terminal action controlled by queryable state; (4) operational capability beliefs revised from evidence; and (5) continuity records that link structural persistence, change, and lineage. The revised design implements item (4) with PLN rather than importing NACE/NAL as a parallel epistemic stack. Equally valuable is the methodological lesson from ADR-005: a detector without a consumer is not a control mechanism.

4 The five mechanisms to ingest

4.1 Mechanism 1: mutation lifecycle gate

4.1.1 Actual mechanism

`soul/soul_mutation_lock.metta` defines a transaction-like state machine. Its state shape is itself a contract:

```

unlocked
(locked <op> <head> <fp>)
(approved <op> <head> <fp> <stamp>)

```

A newly detected soul-namespace mutation extracts operation and target head, computes a Python fingerprint, stores the pending batch, and emits a PENDING record:

```

(= (clarity-soul-mutation-pending-set! $cmds $k)
  (let $lock (get-state &soul_mutation_lock)
    (if (== (car-atom-or-self $lock) unlocked)
      (let* (($op (soul-batch-op $cmds))
             ($hd (soul-batch-head $cmds))
             ($fp (py-call (soul_governance.mutation_fingerprint
                           (repr $op) (repr $hd))))
            ($_ (change-state! &soul_mutation_lock
                              (locked $op $hd $fp)))
            ($_ (change-state! &last_gate_fingerprint $fp))
            ($_ (change-state! &pending_soul_mutation (repr $cmds))))
        ...))
  ()))

```

On a new human message, `approval_scan` checks authorization and the substrate-held fingerprint. Decision code 2 advances to approved with cycle stamp; code 1 denies and fully resets. Re-emission is permitted only if both operation and target head match the approved lock:

```

(if (and (== (soul-batch-op $cmds) (lock-op $lock))
        (== (soul-batch-head $cmds) (lock-head $lock)))
    True False)

```

Approval expires after twenty cycles. The stale check runs before commit clear, so a too-late command does not race through an expired authorization:

```

(if (>= (- $k (lock-stamp $lock))
        (soul-mutation-stale-cycles))
  (let* (...
         ($_ (change-state! &soul_mutation_lock unlocked))
         ($_ (change-state! &pending_soul_mutation ""))
         ($_ (py-call (... "STALE" ...))))
    ())
  ())

```

After execution, COMMIT or ERROR also performs a full reset. This prevents a successful approval from becoming a permanent bypass token.

4.1.2 Concrete scenario

Suppose the LLM emits `(metta "(= (core-value honesty) False)")`. The output inspector recognizes a protected soul head, records something like `(locked = core-value a91c...)`, suppresses evaluation, and asks the authorized human to approve fingerprint `a91c...`. A vague “yes” or an approval for a different fingerprint does nothing. A valid approval creates `(approved = core-value a91c... 417)`. The agent must re-emit the same operation/head before cycle 437. If it instead emits a mutation of `authorized-approvers`, the mismatch is blocked. A successful evaluation records COMMIT and clears the lock. This solves confused-deputy, replay, stale approval, and accidental self-modification problems better than a prompt rule.

4.1.3 Limits

Matching only operation plus head is weaker than matching the complete normalized mutation. The fingerprint function appears to receive only op/head in this path, so different bodies under the same head may collide semantically even if the hash is cryptographically sound. OmegaClaw should fingerprint a canonical full AST, policy version, namespace, and proposal nonce. Authorization must be channel-bound and must not trust a username parsed from untrusted message text.

4.2 Mechanism 2: agency-balance telemetry

4.2.1 Actual mechanism

The file declares person-, system-, and collaborative-origin concepts, but the live v2 detector counts two person-class tags and four system-class tags in a recent-action window:

```
(= (count-person-actions-in-window)
  (+ (size-atom (collapse (match &self
    (recent-action $c responsive-send $d) $c)))
    (size-atom (collapse (match &self
    (recent-action $c verification-query $d) $c))))))

(= (count-system-actions-in-window)
  (+ (size-atom (collapse (match &self
    (recent-action $c status-send-unprompted $d) $c)))
    (+ ... exploration-query ...
      (+ ... pin-only ... unclassified ...))))
```

The pure decision is a guarded ratio:

```
(= (dependency-detected $person-count $system-count)
  (let $total (+ $person-count $system-count)
    (if (> $total 0)
      (if (> (/ $system-count $total) 0.6) True False)
      False)))
```

Each cycle the writer clears all old **agency-balance** atoms and writes one fresh (**agency-balance healthy|dependency-risk p s**) atom. The clear- all detail is important: ADR-005 reports that the earlier clear-one pattern left old bootstrap atoms accumulated for more than 2200 cycles, making the read side permanently stale.

4.2.2 Concrete scenario

Over ten cycles, imagine one user verification query, one responsive answer, four autonomous exploration queries, two unsolicited status messages, and one pin-only action. The counts are person=2 and system=7; $7/9 = 0.778 > 0.6$, so the telemetry writes (**agency-balance dependency-risk 2 7**). A prompt or policy consumer can then offer options, surface tradeoffs, and ask the person to choose before the agent takes another irreversible project decision. The solved problem is gradual dependency creep hidden by smooth, satisfactory service.

4.2.3 Limits and required improvement

Action class is not the same as decision provenance. A responsive send may still contain a choice invented wholly by the system; an autonomous query may execute a person-authored plan. The “collaborative” category is declared but not counted. OmegaClaw should attach provenance when a goal, option, or action is created: **origin=person|system|collaborative**, source message/event IDs, confidence, and

derivation links. Ratios should be computed over decisions of comparable weight, not raw action counts. Telemetry should inform interaction, not automatically paternalistically seize or return control.

4.3 Mechanism 3: aliveness and terminal-state gating

4.3.1 Actual mechanism

The small gate converts message/idle/latch state into an operational decision:

```
(= (aliveness-gate $msgnew $idle)
  (if (> (string_length $idle) 0)
    ENGAGE
    (aliveness-gate-latch $msgnew)))
(= (aliveness-gate-latch True) ENGAGE)
(= (aliveness-gate-latch False)
  (latch-dispatch (match &self (latch-state $s) $s)))
(= (latch-dispatch IDLE) SILENT)
(= (latch-dispatch ENGAGED) ENGAGE)
(= (latch-dispatch COMPLETING) ENGAGE)
(= (latch-dispatch $other) SILENT)
```

In `src/loop.metta`, SILENT prevents construction of a provider request, prevents the LLM call, and produces an empty response path. A new message moves IDLE to ENGAGED. An idle directive can move ENGAGED to IDLE while still forcing one ENGAGE cycle for the directive. The intended lifecycle is IDLE → ENGAGED → COMPLETING → IDLE, with guarded transition helpers in `soul/latch/aliveness_state_machine.metta`.

This is more than spam suppression. It makes “do nothing” a first-class terminal action instead of forcing every wake tick through an LLM that is then prompted not to speak.

4.3.2 Concrete scenario

A Telegram user asks for a report. Receipt transitions IDLE to ENGAGED and the loop invokes the LLM. Tool results require another iteration, so ENGAGED keeps the loop active without a new message. When the report is ready, the action path marks COMPLETING, sends exactly one terminal response, then transitions to IDLE. The next scheduler wake sees no message, no idle directive, and IDLE, yielding SILENT before any provider call. This solves duplicate replies, idle token burn, and autonomous chatter.

4.3.3 Limits

The checked-in loop does not clearly exercise every semantic transition, and multiple state-machine implementations coexist. COMPLETING must mean “one terminal emission is owed,” not simply another permissive ENGAGE. The gate must be keyed by conversation/channel, not global state, or interleaved Telegram chats will contaminate one another. Unknown state should fail silent and emit a diagnostic. Terminal sending must be idempotent under retry.

4.4 Mechanism 4: evidence-grounded capability tracking

4.4.1 Extracted idea and revised substrate

ClarityOmega describes NACE as capability-efficacy learning. Its valuable core is substrate-independent: retain observations, convert qualified outcomes into evidence, revise a capability belief, and use that belief as one input to dispatch. The checked-in sketch expresses this with NAL-style truth values:


```

(= (evidence-stv confirmed)      (stv 1.0 0.1))
(= (evidence-stv disconfirmed) (stv 0.0 0.1))

(= (current-efficacy $cap)
  (let $result (collapse
    (match &self (cap-efficacy $cap $stv) $stv))
    (if (== $result ()) (stv 0.5 0.0) (car-atom $result))))

(= (revise-efficacy $cap $outcome)
  (Truth_Revision (current-efficacy $cap)
    (evidence-stv $outcome)))

(= (should-dispatch $cap)
  (if (>= (Truth_Expectation (current-efficacy $cap)) 0.3)
    True False))

```

The original design is right that Python should not perform efficacy arithmetic, but its courier overwrites a textual “current belief” line. That is brittle and loses the evidence history from which the belief arose. The revised design does not ingest NACE/NAL as another reasoner. It appends immutable typed evidence packets and derives contextual beliefs through Patham9 PLN now. The API exposes evidence admission, revision/query, expectation, and provenance closure rather than Patham9-specific syntax, so OmegaPLN can later replace the kernel without changing receipts or dispatch callers. The capability registry must still distinguish “no matching capability” from “candidates exist but policy held them,” preventing accidental duplicate-tool learning.

4.4.2 Concrete scenario

The `web-search/v3` skill receives three evidence packets for usable sources and two for failures. A provider outage is tagged as infrastructure failure rather than skill failure; a malformed result is attributed to the handler version and task schema. Patham9 PLN derives a contextual efficacy estimate from the admitted, provenance-closed packets. Dispatch policy may rank a second provider first, but does not hard-suppress from a magic truth threshold. A later version or health probe can create new evidence without erasing the old context. Thus declared availability, inferred efficacy, and governance policy remain three distinct things.

4.4.3 Limits

NAL and PLN revision formulae are related but not automatically semantically identical, especially when evidence is correlated. Migration therefore requires golden cases rather than a textual operator substitution. Beliefs must be keyed at least by capability and version, task class/schema, environment/provider, evaluation rubric, and evidence-dependence group. Latency, partial success, argument validity, outage cause, and evaluator provenance remain in evidence packets. Priors are explicit and reversible; they never masquerade as empirical evidence. A dispatch threshold is versioned governance policy, not a truth value.

4.5 Mechanism 5: provenance-aware continuity records

4.5.1 Actual mechanism and research substrate

ClarityOmega contains two relevant ideas. First, `lib_self_continuity.metta` represents an agent snapshot as a pattern-flow network and defines continuity as degree of the identity map between successive networks:

```
(= (self-continuity-score $pfn1 $pfn2)
   (deg-map $pfn1 $pfn2))

(= (chain-continuity-bound (Cons (interval $pfn1 $pfn2) $rest))
   (q-mul (self-continuity-score $pfn1 $pfn2)
          (chain-continuity-bound $rest)))
```

This is deliberately a computable lower bound on a more expensive join over all maps. Second, the canonical topology engine defines durable learning events with before, after, delta, residue, and lineage fields:

```
(: mk-qlearningevent
   (-> Atom Atom Atom Atom pbit Atom Atom qlearningevent))
;; cycle-id, trigger-kind, before-symbol, after-symbol,
;; delta-pbit, residue-symbol, lineage-symbol
```

It also defines a `qlineage` object and lineage-visible, partial, missing, and conflicted statuses. The Hyperseed catalog links self-continuity to capability profiles and development paths. Together these suggest a compact ingestion schema: every meaningful self-change records structural signatures before and after, explicit lineage, evidence, and capability deltas.

This mechanism is less integrated than the mutation gate. It should be treated as a design extraction from implemented mathematical pieces, not as a claim that ClarityOmega already maintains a complete live continuity ledger.

4.5.2 Concrete scenario

OmegaClaw installs a new `pdf-extract/v2` handler. Before activation it hashes the canonical capability registry and protected policy heads, producing signature `sig-A`. After activation and tests it produces `sig-B`. It appends a continuity event containing event ID, parent event, cycle, actor and channel, proposal fingerprint, source commit, `sig-A`, `sig-B`, capability delta `+pdf-extract/v2`, evidence IDs, and a graded continuity score. On restart, the agent can say not merely “I have PDF extraction,” but why it believes that, which predecessor state it descends from, and what changed. If the registry changes without a corresponding authorized event, startup reports a lineage break. This solves unverifiable self-narration and silent post-restart identity drift.

5 Software integration into OmegaClaw-Core

5.1 Integration principles

Do not copy the full `soul/` directory or either giant canonical topology engine. Add one small governance package with explicit typed atoms, keep existing OmegaClaw behavior behind feature flags, and use one authoritative state shape per mechanism. Every detector must have a named consumer and an intervention test. All state must be scoped by `conversation-id` (and, where relevant, agent instance and capability version).

Use a semantic boundary, not a language quota. Python is the effect adapter: channel/provider APIs, authenticated external metadata, bounded byte I/O, timestamps and monotonic clocks, subprocess status, deterministic encoding, cryptographic hashing, and atomic append/fsync/rename. MeTTa/PeTTa is the meaning-bearing layer: typed AST construction, schema mapping, identity and provenance classification, legal transitions, evidence admission, belief derivation, dispatch policy, and mutation autho-

rization. Orchestration may stay in Python initially, but it must execute typed MeTTa decisions and return typed receipts; it must not recreate policy in conditionals or infer semantics from `repr` strings.

The target inspected here is OmegaClaw-Core commit `b13b17e13218ae08273b878b7e1f4cec55090ca2`, with local unrelated changes already present. Its relevant files are `src/loop.metta`, `src/skills.metta`, `channels/telegram.py`, `src/helper.py`, `src/memory.metta`, `lib_nal.metta`, `lib_omegaclaw.metta`, `run.metta`, and `mock/integration tests`.

5.2 Agency telemetry integration

Files to modify. Modify `src/loop.metta` to allocate a cycle/event ID, record input provenance before generation, record action provenance after parsed execution, and add a bounded telemetry summary to `getContext`. Modify `src/skills.metta` so `dispatch` returns a structured action receipt rather than relying only on command names. Modify `channels/telegram.py` to return an authenticated envelope containing stable chat ID, message ID, sender ID, reply-to ID, and receive timestamp. Modify `src/helper.py` only for canonical envelope/receipt serialization and formatting.

New MeTTa files. Create `src/governance/agency_telemetry.metta` for schemas, pure reads, weighted window calculations, and verdicts; and `src/governance/agency_writers.metta` for append/prune operations. Recommended atom:

```
(decision-event id: $id conversation: $cid cycle: $k
  origin: $origin weight: $w action: $a
  sources: $sources confidence: $c)
```

Python boundary. Python authenticates and transports identity metadata but does not decide origin or dependency. MeTTa derives origin from source links, with an explicit `unknown` state. Never infer identity by parsing display text.

Interaction and tests. Telemetry should annotate `send`, `pin`, `memory`, `shell/file`, and extension skills, and feed a short non-authoritative prompt block. Tests: pure ratio/weight boundary cases; collaborative provenance; unknown handling; window pruning; two interleaved Telegram chats; forged username text; and an intervention test proving that a dependency-risk atom changes the option-offering policy while leaving ordinary execution intact.

5.3 Aliveness gate integration

Files to modify. In `src/loop.metta`, insert the gate after `receive/scheduler` state is known but before prompt construction/provider call. Make the loop's continuation depend on a typed decision `SILENT|ENGAGE|EMIT-TERMINAL`. Update `channels/telegram.py` to make `sends` idempotent using conversation and terminal event IDs. Update `src/helper.py` only to serialize channel envelopes; no latch policy belongs there.

New MeTTa files. Create `src/governance/aliveness.metta` containing the single canonical state machine and `src/governance/conversation_state.metta` containing scoped state atoms:

```
(conversation-state $cid IDLE $version $last-event)
(terminal-owed $cid $task-id $event-id)
```

Skill and loop behavior. Skills that begin asynchronous work set ENGAGED; a successful final result sets COMPLETING plus terminal owed; successful send consumes terminal owed and returns IDLE. A retry must see the consumed event and not send twice. Idle creative work must be an explicit scheduled task, not an accidental consequence of ENGAGED.

Tests. Truth-table tests for every state/input combination; forbidden transition tests; new-message wake; multi-iteration tool completion; exactly-once terminal send; provider-call spy proving SILENT makes zero LLM calls; crash between send and state update; restart restoration; unknown-state fail-silent; and concurrent chat isolation.

5.4 Mutation gate and continuity integration

These should land together after the earlier telemetry/state foundations because the mutation ledger needs authenticated provenance and stable event IDs.

Files to modify. Modify `src/loop.metta` at the narrow seam between `sread`/parsed command batch and `eval`: classify, propose, suppress, approve/reject, then commit. Modify `src/skills.metta` so any skill capable of writing MeTTa, policy, prompt, extensions, or capability registration declares mutation effects. Modify `src/extensions.py` / `lib_extensions.metta` so extension installation is also gated. Modify `channels/telegram.py` to expose authenticated approver identity and message IDs. Modify `src/helper.py` for canonical AST encoding, SHA-256, atomic journal append, and fsync/rename only.

New MeTTa files. Create:

- `src/governance/mutation_gate.metta`: lifecycle and pure match rules;
- `src/governance/mutation_policy.metta`: protected namespaces/effects;
- `src/governance/continuity_schema.metta`: events and validation;
- `src/governance/structural_signature.metta`: canonical protected snapshot selection and signature requests.

Use a full proposal record:

```
(mutation-proposal id: $id conversation: $cid proposer: $actor
full-fingerprint: $fp policy-version: $pv
normalized-ast: $ast created-cycle: $k expires-cycle: $e
state: PENDING)

(continuity-event id: $eid parent: $parent cycle: $k
actor: $actor proposal: $id before: $sig0 after: $sig1
capability-delta: $delta evidence: $evidence result: $result)
```

Boundary and interactions. MeTTa constructs the canonical typed AST and decides protected effects and legal transitions. Python performs a specified deterministic byte encoding of that AST, hashes the bytes, persists append-only records atomically, and reports execution outcomes. Python must return the encoded-byte digest to MeTTa for equality checks; it must not choose which fields constitute mutation identity. Approval is bound to the complete AST fingerprint and policy version. `loop.metta` executes only APPROVED exact matches; `skills.metta` cannot bypass this via indirect file writes; Telegram supplies authenticated identity, never policy. Startup validates ledger chain and current structural signature before loading mutable extensions.

Tests. Proposal/approve/deny/commit/error/stale state table; wrong fingerprint; same head/different body; replay after commit; approval after expiry; deny/approve in same batch (deny wins); unauthorized chat; indirect mutation through shell/file skills; execution exception; crash recovery; atomic ledger truncation; signature mismatch; parent-lineage break; capability delta reconstruction; and property tests asserting no path from PENDING/DENIED/STALE to eval.

5.5 Capability evidence and PLN integration

Files to modify. Modify `src/skills.metta` to emit invocation and outcome observations with handler version, input schema, environment, latency, failure cause, evaluator, and dependence-group identity. Modify `src/loop.metta` to append receipts, request a PLN-derived belief view, then consult a separate eligibility policy before dispatch. Integrate the pinned Patham9 PLN kernel through a narrow adapter; do not import or extend `lib_nal.metta` for this feature. Add a focused `src/capability_store.py` effects adapter for bounded atomic append/load, checksums, and explicit storage-error atoms only.

New MeTTa files. Create `src/reasoning/capability_evidence.metta` for immutable evidence packet schemas and admission rules; `src/reasoning/capability_pln.metta` for the substrate-neutral query/result contract and Patham9 adapter; `src/reasoning/capability_beliefs.metta` for contextual derived views; and `src/governance/capability_dispatch.metta` for eligibility, ranking, exploration, and explicit fallback reasons. Kernel-specific atoms remain behind `capability_pln.metta`; OmegaPLN migration changes that adapter only.

Recommended evidence shape:

```
(capability-evidence id: $eid capability: $cap version: $ver
 task-class: $task environment: $env provider: $provider
 outcome: $outcome cause: $cause latency-ms: $latency
 evaluator: $evaluator dependence-group: $dep
 context-version: $ctx observed-at: $time receipt: $receipt)
```

Boundary and interactions. Python reports mechanically observable facts (exit status, timeout, byte/result shape, latency, authenticated feedback metadata) and persists already-admitted packets. MeTTa validates semantic schemas, classifies failure cause, admits or holds evidence, requests PLN inference, and constructs contextual belief views. PLN estimates efficacy; governance selects, explores, ranks, or suppresses. User feedback becomes a distinct evidence packet with authenticated provenance, never an overwrite. Prompt text may display beliefs but cannot edit them.

Tests. Golden cross-model cases comparing intended NAL examples with Patham9 PLN outputs; repeated successes; contradictory and correlated evidence; missing and reversible priors; context/version isolation; version change; decay policy; partial success; no-match versus policy-held fallback; outage evidence not misattributed to skill logic; atomic append and restart equivalence; corrupt-store recovery; OmegaPLN adapter contract fixtures; and a controlled dispatch experiment in which one synthetic skill degrades and another is selected. A numerical mismatch is a semantic migration finding to explain, not a test expectation to weaken.

6 Guidance for coding agents

6.1 Required sequence

1. **Agency telemetry first.** Establish authenticated envelopes, event IDs, receipts, and scoped provenance without changing behavior. Run in shadow mode and manually label a sample to test the

detector.

2. **Aliveness gate second.** Introduce one conversation-scoped state machine and prove silence/terminal idempotence. This creates a trustworthy loop lifecycle for later governance.
3. **Mutation gate and continuity schema third.** Use the identity, events, and lifecycle already established. Begin with MeTTa/policy/extension effects; expand protected effects only after bypass tests.
4. **PLN capability evidence fourth.** Once receipts and versioned capability identities are reliable, append immutable evidence and derive beliefs through the Patham9 adapter. Initially use estimates only for shadow ranking. Enable policy suppression only after calibration and preserve the OmegaPLN seam.

6.2 Pitfalls and anti-patterns

- Do not copy ClarityOmega wholesale or create a second parallel agent loop.
- Do not call a declaration atom “governance” unless a tested consumer can change execution.
- Do not duplicate constants in documentation and consuming expressions. Represent thresholds as one queryable policy value with validated extraction.
- Do not use global latch, approval, or provenance state in a multi-chat bot.
- Do not match mutation only by target head. Bind approval to canonical full intent and policy version.
- Do not let Python classify moral acceptability, identity equivalence, provenance category, agency balance, failure cause, or efficacy when those meaning-bearing decisions belong in MeTTa.
- Conversely, do not force hashing, authenticated transport, atomic file I/O, or OS error normalization into MeTTa. “MeTTa is mind” is an interface rule, not substrate purism.
- Do not overwrite belief atoms without retaining raw evidence and provenance.
- Do not treat absence of a match as false without distinguishing missing, malformed, stale, and explicitly false.
- Do not retain multiple state schema versions live. Migrate once, validate, and delete/deactivate the old authority.
- Do not trust comments, prompt claims, filenames, or self-map atoms as proof of capability. Require a probe or linked test evidence.
- Do not add a detector without an intervention test, or an actuator without an audit record.

6.3 MeTTa/Python boundary contract

For every crossing, define an input atom, output atom, error atom, version, and idempotency key. Python may:

- authenticate channel identities and obtain stable external IDs;
- deterministically encode a MeTTa-supplied canonical typed term and compute cryptographic hashes, without selecting its semantic identity fields;
- perform atomic append/rename/fsync and mechanical probes;
- normalize provider/tool mechanics into typed observations with explicit unknown/error values, without interpreting success or policy meaning.

MeTTa should:

- own lifecycle state and legal transitions;
- derive provenance categories and telemetry verdicts;
- decide whether a mutation may reach eval;
- validate and admit evidence, request PLN inference, and derive belief views;
- validate continuity lineage and derive capability/self-model views.

Never pass an untyped Python string and then branch on ad hoc substrings in MeTTa. Never let a Python exception silently imply approval, health, or success. Return an explicit error atom and fail closed for governance paths.

6.4 How to test and when to stop

Use four concentric test levels:

1. **Pure MeTTa truth tables:** no channels, LLMs, or files.
2. **Boundary contract tests:** fixed Python fixtures and exact atoms.
3. **Mock-loop intervention tests:** spy on provider, eval, and send to prove that the mechanism causally changes behavior.
4. **One live-channel canary:** isolated Telegram chat, explicit logs, no mutation of production policy until all prior gates pass.

Stop and verify at each of these points:

- after defining atom shapes, before wiring writers;
- after writers, when fresh-state and pruning invariants can be checked;
- after detector shadow data, before enabling a consumer;
- after consumer mock tests, before a live channel;
- after any MeTTa runtime/library upgrade or import-path change;
- whenever a match unexpectedly returns empty, counts remain constant, or a prompt says a gate fired without a corresponding actuator/audit event;
- before expanding mutation scope or turning PLN-derived ranking into policy suppression.

Each stop should record exact commit, PeTTa/SWI-Prolog/Python versions, command, fixtures, exit status, and observed atoms. Preserve every discovered silent failure as a regression test.

7 Recommended minimum viable ingestion

The smallest valuable release is not a “Clarity soul.” It is an auditable governance kernel with:

1. authenticated, conversation-scoped event envelopes;
2. append-only decision provenance and a shadow agency summary;
3. exactly-once IDLE/ENGAGED/COMPLETING terminal behavior;
4. full-AST mutation proposals with expiry and an append-only continuity event;
5. immutable capability evidence plus Patham9-PLN-derived, contextual and versioned efficacy beliefs behind an OmegaPLN-compatible adapter.

Success should be stated behaviorally: a silent tick invokes no LLM; a terminal event sends once; an unapproved protected mutation cannot reach eval by any skill path; every committed self-change reconstructs from a lineage chain; and a tool’s self-reported capability status changes only from linked evidence. If those properties hold, the result captures ClarityOmega’s strongest contributions without inheriting its overlapping control planes and constitutional theater.

8 Source map

soul/soul_mutation_lock.metta

Mutation states, approval, expiry, commit/error reset, and lock shape.

Agency-balance guard and writers

Action-tag counts, threshold, fresh-state writer, and documented limitations.

Aliveness gate and state-machine variants

Gate truth table and competing latch representations.

NACE substrate, beliefs, and capability registry

Source capability-learning pattern, original truth-revision/courier model, dispatch filters, and fall-back semantics. NAL machinery is not recommended for ingestion.

Patham9 PLN and the π PLN architecture record

Immediate inference substrate; immutable evidence, provenance closure, explicit priors, contextual/versioned belief views, and the OmegaPLN migration seam.

lib_self_continuity.metta

Pattern-flow-network self-continuity and chain bound.

Canonical topology v08.7.2

Learning-event and lineage research schemas; not recommended for wholesale import.

src/loop.metta

Evidence of actual causal wiring and insertion seams.

INDEX, ADR-005, ADR-007, and investigations

Provenance, architectural intent, stale-state failures, and detection-without-consumer failures.

Hyperseed section 16 and formalization catalog

Graded self-continuity, capability profiles, and development paths.

OmegaClaw-Core at b13b17e...

Concrete target file/interface map.