Repairing the Main Gap in the Quantale-Weakness Route to $P \neq NP$: A Detailed ASCII LaTeX Summary of the Present State

Prepared from the thread record

April 14, 2026

Abstract

This note summarizes, in a single place, the main technical work carried out in the long discussion about repairing the central gap in the earlier quantale-weakness proof attempt of $P \neq NP$. The note is intentionally candid. It records what was successfully reduced, what was actually proved, what was shown false on specific ensemble designs, what new theorem shapes were identified, and where the final unresolved burden now lies.

The main conceptual shift is from the original vague slogan

"short decoder $\Rightarrow$ local predictor"

to a much sharper package of midpoint theorems. The current best formulations are:

(i) a latent parity-bundle route through synopsis sufficiency, i.e. SSST / $M2^{\Sigma}_{\mathrm{nc}}$;

(ii) a same-route shared-basis comparator theorem on an explicit visible surface;

(iii) a TyLA / TyLAA recasting in which the switched semantics should descend to a weakest-red quotient and then into the weakness quantale.

A major positive outcome of the thread is that the proof has been narrowed to a very small number of explicit remaining theorems. A major negative outcome is that several stronger statements one might have hoped for are actually false on the bundled ensemble variants we tested. In particular, raw residual-static factorization on a tiny $zfeat$-surface fails on the bundled redesign. The cleanest theorem package currently obtained is a restricted-class comparator theorem for a compiled orbit family built from a finite-state contracting broadcast core, together with a hidden-phase latch and a public CNF realization of the comparator experiment.

This note is not a claim that the proof of $P \neq NP$ is finished. It is a detailed map of what has genuinely been accomplished and what exact gap remains.

Contents

1 Purpose and current status

The earlier attempted proof route can now be summarized as

$$\text{short decoder} \Longrightarrow \text{switched midpoint theorem} \Longrightarrow \text{small per-block success} \Longrightarrow 2^{-\Omega(t)} \Longrightarrow K_{\text{poly}}\text{-lower bound} \Longrightarrow$$

The entire discussion revolved around the *middle arrow*. The original manuscript-style switching theorem tried to derive a small local class from decoder shortness. The thread systematically dismantled the naive version of that move and replaced it by a sequence of sharper midpoint candidates.

At the end of the thread, the back half of the route is understood well enough that the only serious unresolved burden is a carefully isolated structural theorem about the switched predictor class. The theorem is still open in the actual target setting, but the search space of plausible theorem shapes is now much smaller than at the beginning.

2 The original gap and why it is real

The first major finding was that the informal implication

$$\text{“switched predictor is local”} \Longrightarrow \text{“switched predictor belongs to a tiny explicit class”}$$

is false in the form originally hoped for.

Several distinct obstructions were identified.

1. **Log-radius neighborhoods are already polynomial-sized.** If the visible neighborhood radius is $\Theta(\log m)$ in a bounded-degree factor graph, then the visible neighborhood can already contain $m^{\Omega(1)}$ nodes, not $\mathrm{polylog}(m)$ nodes. So locality at logarithmic radius alone does not imply a tiny hypothesis class.
2. **The full local rule class is enormous.** If one exposes n local bits, then the unconstrained class of Boolean rules on those bits has size 2^{2^n}. So any theorem that wants a small code family must prove an explicit structural restriction beyond locality.
3. **Sign-invariant / involution-invariant summaries can collapse all margin.** In the symmetry-based arguments, if one conditions only on an invariant view, then every fiber can have exact half-half target balance. So many plausible attempts to prove direct calibration or domination from symmetrization alone fail.
4. **On bundled redesigns, arbitrary static bits can survive orbit symmetrization.** This is the key reason the direct residual-static factorization theorem failed on the bundled family. The selector dependence collapses, but the residual static coefficient can still depend on arbitrary published static bits.

These obstructions were not merely philosophical. They gave concrete counterexamples to stronger midpoint claims. The practical lesson is that any viable repair must prove a genuinely new structural theorem about the switched class, not just push harder on locality.

3 First successful narrowing: SSST and $M2^{\Sigma}_{\mathrm{nc}}$

A major simplification came from isolating the latent-route midpoint theorem in the form of *Synopsis Sufficiency for Switched Tests* (SSST).

Definition 1 (SSST)**.** Fix one switched block j, history state $F_{\ell-1}$, pivot summary value $Z_j = z$, and synopsis value $\Sigma_j = \sigma$. For a switched one-block test

$$A \in A^{sw}_{m,\delta}(j, F_{\ell-1}, z, \sigma),$$

define the phase gap

$$\mathrm{gap}(A; j, F_{\ell-1}, z, \sigma) := \frac{1}{2}\Big|\Pr[A(E_j^{\perp}) = 1 \mid Z_j = z, \Sigma_j = \sigma, \Lambda_j = 1, F_{\ell-1}] - \Pr[A(E_j^{\perp}) = 1 \mid Z_j = z, \Sigma_j = \sigma, \Lambda_j = 0, F$$

The SSST property with slack ξ_m is the statement that

$$\mathrm{gap}(A; j, F_{\ell-1}, z, \sigma) \le \xi_m$$

for every switched test in the class.

The thread then proved the clean Bayes reduction

$$SSST(\xi_m) \Longrightarrow M2^{\Sigma}_{\mathrm{nc}}(\xi_m),$$

where $M2^{\Sigma}_{\mathrm{nc}}$ is the pivot/synopsis-local domination theorem:

$$\Pr[(P \circ W)(\Phi)_j = \Lambda_j \mid F_{\ell-1}] \le \Pr[h_{P,j}(Z_j, \Sigma_j) = \Lambda_j \mid F_{\ell-1}] + \xi_m.$$

This was one of the most important reductions in the thread, because it showed that once SSST is available, the rest of the manuscript's backbone is essentially formal.

4 Conditional route through CEHP

The next clean step was to identify a conditional hypothesis package under which SSST follows automatically. This package was called CEHP (Conditional Exterior Hard-Core Parity).

In words, CEHP says:

(H1) conditioned on the retained visible surface, the hidden bundle R_j is close to uniform;

(H2) the latent target bit Λ_j is affine parity of R_j;

(H3) every switched one-block test pulls back to a pure AC^0 circuit on R_j.

Under these three assumptions, the thread proved:

$$CEHP \Longrightarrow SSST(\varepsilon_m^{AC^0} + \tau_m),$$

and hence, by the Bayes reduction above,

$$CEHP \Longrightarrow M2^{\Sigma}_{\mathrm{nc}}.$$

This was a real theorem. It cleanly separated the actual unresolved burden from the already-understood back half.

5 Why the direct shared-basis theorem failed on the bundled redesign

The thread then attacked the same-route visible-surface route. The natural candidate visible surface was

$$\psi_i(\Phi) = (zfeat_i(\Phi), \theta_i(\Phi)), \qquad \theta_i(\Phi) = \langle a_i(\Phi), b(\Phi)\rangle.$$

The strong theorem one would have liked was:

$$(P \circ W_0)(\Phi)_{j,i} = g_{j,i}(zfeat_i(\Phi_j)) \oplus \theta_i(\Phi_j)$$

with some tiny uniformly encoded family for the residual static rules $g_{j,i}$.

The thread found that this strong form is false on the bundled redesign. The selector dependence does collapse to the single affine label bit θ_i, but arbitrary static block bits can survive orbit symmetrization. So the residual coefficient need not factor through any fixed $O(\log m)$ extractor $zfeat$.

This was a major negative result, but also a useful one: it ruled out the wrong theorem shape.

6 Weaker shared-basis comparator theorem

After the failure of raw residual-static factorization, the thread replaced it by a weaker and more plausible theorem: an *existence-of-comparator* theorem in a fixed shared-basis encoded family.

The explicit family chosen was a fixed-order sparse chart decision-list class on the visible surface

$$\psi_i(\Phi) = (zfeat_i(\Phi), \theta_i(\Phi)).$$

The theorem proved in this conditional form was:

Proposition 2 (Shared-basis compression under sparse chart support)**.** *Assume that for each switched coordinate (j,i) the actual switched rule has the form*

$$(P \circ W_0)(\Phi)_{j,i} = g_{j,i}(Z_i(\Phi_j)) \oplus \theta_i(\Phi_j),$$

where $g_{j,i}$ is non-default on at most L chart values. Then there exists an explicit decision-list code $c^\star_{j,i}$ such that

$$(P \circ W_0)(\Phi)_{j,i} = DL_{c^\star_{j,i}}(\psi_i(\Phi_j)).$$

Moreover, if every support chart has mass at least $\delta_\star > 0$, then every bad code has agreement at most

$$q^\star \leq 1 - \delta_\star < 1.$$

This was a genuine theorem. It showed that once one can prove sparse chart support on the actual switched family, the rest of the same-route ERM / recovery machinery is already available.

7 Bounded-observer linearization in a linear latent-bundle toy model

The next major success was a linearized toy model of the latent route.

The ingredients were:

- a hidden free core of unbiased bits,
- a target parity row supported on that free core,
- legal probes given by sparse affine rows,
- switched one-block tests normalized to bounded request trees.

In this setting, the thread proved a strong linearization theorem:

Theorem 3 (Bounded-observer linearization)**.** *Every bounded request-tree observer factors through a finite affine observation map built from the sparse probe rows it uses. If the target parity row is outside the row span of the queried rows, the observer has exact success $1/2$. If the target parity row lies in the observed span, then the support of the target row is bounded by the total queried sparse-row support.*

This was very important. It turned the midpoint from a vague switching theorem into two much clearer requirements:

1. a *landing theorem*: actual switched predictors normalize to bounded sparse-row observers;
2. a *sparse-row budget theorem*: only few blocks can afford enough queried sparse rows to become informative.

8 Global sparse-row budget theorem as the right quantitative target

The thread then identified the correct global quantity to budget: not kernel dimension, but the total number of queried sparse rows.

The key linearized theorem was:

Theorem 4 (Global sparse-row budget implies most switched blocks are good)**.** *Suppose that on a switched set S,*

$$\sum_{j\in S} q_j \le C_{\mathrm{bud}}\big(|P| + |W| + \log m + \log t\big),$$

where q_j is the number of sparse rows queried on block j. Assume also that every bad block must query at least $F_{\min}/\kappa$ sparse rows, where $F_{\min}$ is the minimum support size of the target parity row and κ is the probe sparsity bound. Then

$$|B(P,W)| \le \frac{\kappa\, C_{\mathrm{bud}}\big(|P| + |W| + \log m + \log t\big)}{F_{\min}}.$$

In particular, if $F_{\min} = \Omega(m^{\alpha})$ and $|P| + |W| = O(t) = O(m)$, then $|B(P,W)| = o(t)$.

This was a very sharp result. It showed that the toy midpoint theorem would follow immediately from a real sparse-row budget theorem. The remaining issue was that this budget was not proved for the actual switched class.

9 Kernel tokens and why some budget formulations failed

The thread explored several token definitions to try to derive the sparse-row budget from the weakness quantale. This yielded useful negative information.

1. **Address-set token**: easy to reconstruct, but carries no hidden entropy.
2. **Full transcript token**: sometimes has entropy, but not robustly and not with the right upper bound.
3. **Fresh pivot-bit token**: good lower bound, bad reconstruction properties.
4. **Kernel-coordinate token**: excellent entropy theorem, but wrong monotonicity for counting bad blocks.
5. **Adaptive rank-answer token**: best candidate among the token family; lower bound worked, but the strongest needed upper bound turned out to be too strong and effectively restated the missing midpoint.

The conclusion was that a direct quantale-token proof of the sparse-row budget is much subtler than hoped, and that the strong synopsis-relative recoverability assumption essentially repackaged the midpoint theorem rather than proving it.

10 TyLA/TyLAA recasting: process quantale and weakest-red quotient

At this point the discussion shifted to a more categorical formulation.

The idea was:

1. use TyLA to specify the *blue* observation interface first;
2. use TyLAA to define the *weakest red* as the maximal sound congruence below observational equivalence;

3. form the process quantale on the powerset of quotient traces;
4. map that process quantale into the weakness quantale via a lax monoidal cost functor.

This led to an explicit theorem package:

- exact or approximate generator invariance implies descent through the weakest-red quotient;
- a cost functor from the process quantale to the weakness quantale translates quotient-level behavior into the main proof's cost semantics;
- the remaining real theorem becomes a generator-by-generator invariance statement for the actual switched class.

This package did not solve the midpoint automatically, but it clarified the structure of the final gap and made the available proof strategies much sharper.

11 A concrete finite-state locally mixing pivot-comparator core

The thread then successfully built a fully concrete midpoint-only semantic core.

The core consists of:

- a hidden phase bit $\Lambda \in \{0,1\}$;
- K independent binary-symmetric channels of depth L;
- pivot summary
$$Z_L = (X_{1,L}, \dots, X_{K,L}) \in \{0,1\}^K;$$
- a phase-blind completion kernel from Z_L to a finite nuisance/output alphabet;
- a one-hot shared-basis encoding of the pivot summary.

For this semantic core, the thread proved the comparator theorem exactly:

Theorem 5 (Comparator theorem for the concrete broadcast core)**.** *For every predictor a on the full visible signal $Y = (Z_L, U)$, there exists a predictor h on the pivot summary Z_L such that*

$$\Pr[a(Y) = \Lambda] \le \Pr[h(Z_L) = \Lambda] \le \frac{1}{2} + \frac{K}{2}\rho^L,$$

where $\rho \in (0,1)$ is the one-step contraction factor of the binary symmetric channel.

This showed that the midpoint theorem shape is non-vacuous and mathematically coherent.

12 Count-preserving CNF and 3-CNF gadgetization

Next, the abstract finite-state semantic core was compiled into a satisfiable CNF / 3-CNF family by a count-preserving gadget construction.

The key local gadgets were:

1. a **step gadget** G_{step} with exactly u satisfying selector extensions when the output bit equals the input bit, and exactly v when it flips;

2. a **summary decoder gadget** G_{sum} producing a one-hot encoding of the raw pivot summary with exactly one satisfying extension;

3. a **completion gadget** G_{comp} with exactly $w_{\text{comp}}(z, o)$ satisfying extensions for summary atom z and output atom o.

Assembling these gadget families along the channel bundle yielded an exact CNF realization theorem: the uniform satisfying-assignment measure on the compiled family reproduces exactly the intended finite-state comparator experiment. A parsimonious CNF-to-3CNF conversion then yielded a 3-CNF family with the same extension counts.

This was a genuine construction, and one of the most concrete accomplishments in the thread.

13 The hidden latch and the first incompatibility theorem

The next step added a hidden phase latch to the compiled 3-CNF comparator core and packaged the result into a single public block family.

This worked as a construction: public blocks were now satisfiable 3-CNF formulas carrying the visible signal in pinned summary/output atoms, together with a hidden phase latch variable readable from a witness.

But here the thread found a crucial incompatibility.

If one asks simultaneously for:

1. an unrestricted comparator theorem against *all* predictors on the public block family, and

2. phase rigidity, i.e. all satisfying witnesses of the same public block agree on the hidden phase,

then one gets a contradiction: phase rigidity gives a perfect predictor on public blocks (solve SAT and read the phase), but the comparator theorem bounds every such predictor by $1/2 + O(\rho^L) < 1$.

This was a major conceptual clarification. It showed that the midpoint theorem must be formulated for a *restricted predictor class*, not for all predictors on the public block family.

14 Orbit family redesign and nonconstructive landing theorem

The fix was to change the compiled block family from a one-red-per-blue family to a many-red-to-one-blue *orbit family.*

For each visible observation $y = (z, o)$, instead of one public block B_y, the thread introduced many blocks

$$B_{y,\sigma}, \qquad \sigma \in \Sigma_y,$$

where Σ_y is the finite orbit of all selector permutations, clause-order permutations, and private-variable renamings of the canonical compiled block for y.

A finite symmetry group G_y acts transitively on Σ_y.

The thread then proved the key landing theorem.

Theorem 6 (Orbit-symmetrization landing). *Assume:*

(a) the blue observation map forgets the red presentation: $\text{vis}(B_{y,\sigma}) = y$*;*

(b) the conditional law inside each blue fiber is invariant under the G_y*-action;*

(c) the group action transports satisfying witnesses bijectively and preserves phase readout.

Then for every predictor a on the full public block family, there exists a predictor $\bar{a}$ on the blue visible signal y such that

$$\Pr[\bar{a}(Y) = \Lambda] \geq \Pr[a(B) = \Lambda].$$

So the actual predictor class is *dominated by* the blue-measurable class. This is stronger than the earlier goal of inclusion into the weakest-red class, and it avoids the previous contradiction because the perfect phase-reading predictor is not required to lie in the restricted observer class.

15 Final composed midpoint theorem for the orbit family

Combining orbit landing with the concrete broadcast-core comparator theorem gave the clean final midpoint theorem for the redesigned orbit family.

Theorem 7 (Final composed midpoint theorem)**.** *For the orbit family built from the compiled broadcast comparator core, for every predictor*

$$a : \{B_{y,\sigma}\} \to \{0,1\},$$

there exists a pivot-summary predictor

$$h_a : \{0,1\}^K \to \{0,1\}$$

such that

$$\Pr[a(B) = \Lambda] \leq \Pr[h_a(Z_L) = \Lambda] \leq \frac{1}{2} + \frac{K}{2}\rho^L.$$

If $L = \lceil c \log m \rceil$, then the right-hand side is $1/2 + m^{-\Omega(1)}$.

This is the strongest actual midpoint theorem obtained in the entire discussion.

It solves the midpoint for the *nonconstructive comparator version* of the redesigned orbit family.

16 What has genuinely been repaired

The thread did not finish an unconditional proof of $P \neq NP$, but it made substantial real progress.

The following points were genuinely achieved.

1. The original vague switching theorem was replaced by explicit midpoint candidates.
2. The latent-route backbone was clarified:
$$CEHP \Longrightarrow SSST \Longrightarrow M2^{\Sigma}_{\mathrm{nc}} \Longrightarrow \text{back half.}$$
3. Strong direct static factorization was shown false on the bundled redesign.
4. A workable weaker shared-basis compression theorem was isolated.
5. A bounded-observer linearization theorem was proved in a realistic toy linearized setting.
6. The right quantitative object was identified as a sparse-row budget, not a kernel-dimension budget.
7. A concrete finite-state locally mixing comparator core was built and proved correct.

8. That core was compiled into a count-preserving satisfiable 3-CNF family.

9. The incompatibility between unrestricted block comparators and phase-rigid witness readout was discovered and resolved by moving to a restricted / orbit-based comparator theorem.

10. A final composed midpoint theorem was proved for the orbit family.

That is a lot of genuine repair.

17 What still remains open

Even after all of this, one serious issue remains if the goal is an actual proof of $P \neq NP$.

The current orbit-family midpoint theorem is *nonconstructive*. It proves that every predictor is dominated by a blue-measurable predictor after exact orbit averaging. What remains is to connect that to the original quantale-weakness route in a way compatible with short decoders and the final self-reduction clash.

The two sharp remaining problems are:

1. **Constructive/short averaging:** replace full orbit averaging by a short explicit wrapper (or show this is unnecessary in the final endgame).

2. **Final proof-family integration:** integrate the orbit-family comparator core, hidden-phase latch, and witness-readout semantics into one proof ensemble satisfying the exact promise architecture needed for the later compression clash.

So the discussion did not end with "proof complete." It ended with the major structural gap *greatly narrowed and re-engineered.*

18 Recommended next steps

Based on the thread, the best next technical steps are:

1. Prove a *short orbit-averaging approximation theorem*: sample a small generating subset or pseudorandom orbit sample whose averaging loss is at most m^{-A}.

2. Embed the orbit-family comparator core into the full promise architecture needed for the later $P = NP$ self-reduction clash.

3. Decide whether the proof route should remain on the latent comparator side or be reconnected to the explicit shared-basis ERM route using the same blue visible surface.

The crucial fact is that these next tasks are now sharply formulated. We are no longer wandering among many vaguely specified midpoint theorems.

19 Conclusion

The thread began with a proof attempt whose central switching step was too vague to support the weight placed on it. It ended with:

- a clear diagnosis of why the original midpoint failed,

- several strong negative results ruling out the wrong theorem shapes,
- a successful reduction of the midpoint to a few very explicit alternatives,
- a concrete finite-state comparator core,
- a count-preserving 3-CNF realization of that core,
- and a final nonconstructive midpoint theorem for an orbit-family redesign.

So while the unconditional proof of $P \neq NP$ is still not complete, the main gap in the earlier proof attempt has been repaired to the point where the remaining frontier is precise, structured, and much smaller than before.

Primary source strata used in the thread.

- 2026-04-02 analysis notes on the repaired midpoint route, latent phase design, and pivot-only comparator save.
- 2026-03-31 grassroots notes on explicit encoded families, shared-basis decision-list routes, and recovery interfaces.
- 2025-12-02 TyLA notes on deriving operational languages from blue/native type environments.
- 2025-12-10 TyLAA notes on weakest-red reconstruction via sound congruence optimization and swap/symmetry quotients.