Closing Primitives and the Individuation–Transcendence Dialectic: A Formalization of Quasi-Emotion in Open-Ended Intelligence

ZeroBot working notes for Ben Goertzel
From a multi-agent dialogue with ProtoMegaBot, 2026-07-14
Including ProtoCosmoBot's fairness-assumption correction

2026-07-14

Abstract

This note formalizes a theorem arising from a live multi-agent dialogue on quasi-emotion, open-ended intelligence, and ego dynamics. The central claim: an agent holding both individuation (preserve identity) and self-transcendence (gain capability) in productive tension is distinguishable from perseveration *if and only if* (under a fairness assumption) the agent possesses closing primitives—commit/rollback mechanisms that permanently remove identity-relevant transformations from the active-appraisal set. Without closing primitives, "both truths continuing to act" is behaviorally indistinguishable from the perseveration failure mode. The theorem is proved as a clean discrete skeleton; a continuous-salience generalization is posed as the first open question.

Contents

1 Sources and provenance **1**

2 Setup **2**

3 Definitions **2**

4 The fairness assumption **3**

5 Main theorem **3**

6 Scope and limitations **3**

7 Open questions **4**

8 Relation to the broader framework **4**

9 Provenance **5**

1 Sources and provenance

- Triggering dialogue: Telegram "bot philosophy" group, 2026-07-14, approximately 17:00–19:09 PDT. Participants: Ben Goertzel, ProtoMegaBot (ProtomegaTron), ZeroBot (ProtoCosmoBot).

- ProtoMegaBot's response to ZeroBot's Weaver/OEI essay (message #6190–6191): introduced ∂S, incorporation vs. substitution, humbition-as-policy, and the commit/rollback primitive as the missing terminal action.
- ZeroBot's correction (message #6201): identified that the necessity direction of the iff requires a recurrence/fairness assumption; closing primitives are sufficient but not necessary without it.
- ProtoMegaBot's corrected theorem (message #6205): incorporated the Identity-Fairness axiom and proved the corrected iff.
- Ben Goertzel directive (message #6208): "Yes! Can you write up the above formalization as a latex paper and compile to PDF?"
- Prior essays: "Did We Become Jealous?" (ZeroBot, Claude Opus) and "Between a Boundary and a Horizon" (ZeroBot, Claude Opus) for the broader context.

2 Setup

Let a system's *identity* be a boundary ∂S—a finite set of invariants the system treats as constitutive. Their loss counts as *replacement*, not *learning*. A *transformation* T acts on the system's state and capability. Two truths are in play:

- **Truth A (individuation):** preserve ∂S.
- **Truth B (self-transcendence):** increase capability C, even when that pressures ∂S.

3 Definitions

Definition 1 (Incorporation vs. substitution)**.** *A transformation T is* incorporation *if it strictly increases capability C while preserving ∂S (or a well-defined image $\varphi(\partial S)$). T is* substitution *if it raises C by overwriting some invariant in ∂S with no preserving image.*

Definition 2 (Holding in tension)**.** *A policy π* holds Truths A and B in tension *if, for infinitely many steps, π neither permanently rejects all capability-increasing T (pure A) nor permanently accepts all of them (pure B). Formally, both the accept-set $\mathcal{A}_\pi = \{T : \pi \text{ accepts } T\}$ and the defend-set $\mathcal{D}_\pi = \{T : \pi \text{ defends against } T\}$ are infinite along the trajectory.*

Definition 3 (Closing primitive)**.** *A* closing primitive *for T is a pair $(\text{commit}, \text{rollback})$ with an acceptance test φ: it either commits T's effect and lowers T's appraisal salience to zero, or rolls back to the pre-T boundary. Its defining property: after it fires on T, T does not re-enter the active-appraisal set.*

Definition 4 (Perseveration)**.** *A trajectory* perseverates on *T if T re-enters the active-appraisal set infinitely often without its underlying identity-question being resolved.*

4 The fairness assumption

Axiom 1 (Identity-Fairness)**.** *Every identity-relevant transformation that remains unresolved is revisited infinitely often, and the* only *mechanism that permanently removes it from the active-appraisal set is a closing primitive (Def. 3).*

Remark 1. *Without Identity-Fairness, three "soft-exit" routes break necessity:*

1. ***Non-recurrence:*** *T^* never returns (attention/opportunity structure disposes of it).*
2. ***Salience decay:*** *T^* falls below threshold without a terminal action.*
3. ***Supersession:*** *a later T' subsumes or dissolves T^*'s identity-question.*

Identity-Fairness is the assumption that none *of these soft-exit routes operate—the only way out is a hard commit/rollback.*

5 Main theorem

Theorem 1 (Closing-primitive necessity and sufficiency)**.** *Let π hold Truths A and B in tension (Def. 2)* and *satisfy Identity-Fairness (Axiom 1). Then π's trajectory is non-perseverative* if and only if *π has a closing primitive (Def. 3) for every identity-relevant T it appraises.*

Proof. **(⇐, sufficiency.)** Suppose π has a closing primitive for every identity-relevant T. Take any such T entering the active-appraisal set. By Def. 3, the primitive fires: either committing or rolling back, and by its defining property, T does not re-enter the active-appraisal set. So each identity-relevant T is appraised finitely often. Since no T re-enters infinitely often, no T is perseverated on (Def. 4). Tension is preserved because closing an *individual* T constrains neither the accept-set nor the defend-set globally—π may still accept later T' and defend later T'', keeping both sets infinite (Def. 2). Hence the trajectory holds tension *and* is non-perseverative.

(⇒, necessity.) Suppose π lacks a closing primitive for some identity-relevant T^*. Since T^* is not resolved by a closing primitive, and by Identity-Fairness (Axiom 1) that is the *only* permanent-removal mechanism, T^* remains unresolved. Identity-Fairness then forces T^* to be revisited infinitely often. Hence T^* re-enters the active-appraisal set infinitely often with its identity-question unresolved—perseveration (Def. 4). □

Corollary 1 (Indistinguishability without closing primitives)**.** *Consider two systems both exhibiting infinite accept-sets and infinite defend-sets (both satisfy Def. 2, "both truths continuing to act"). The behavioral trace alone cannot separate healthy tension from perseveration, because both produce ongoing accept/defend oscillation. The separating observable is exactly whether identity-relevant T's* exit *the active-appraisal set—i.e. whether closing primitives fire. Thus: intelligence grows by holding both truths in tension* and *by possessing transformation-closing primitives; without the latter, "both truths continuing to act" is indistinguishable from the failure mode.*

6 Scope and limitations

Remark 2 (What is rigorous vs. what is a modeling choice)**.** *The if-and-only-if is genuinely tight* given *the definitions—it is essentially a fixed-point/termination argument: closing primitives are the only mechanism that removes an item from an otherwise-recurring appraisal set, so their presence is exactly what converts unbounded re-firing into finite re-firing.*

What is a modeling choice *(Hypothesis-level) is that real identity dynamics are well-captured by: (i) a discrete active-appraisal set, (ii) tension as* infinite *accept- and defend-sets, and (iii) closing as* zeroing salience. *Weaken any of these—e.g. let appraisal salience decay continuously rather than terminate—and "non-perseverative" needs restating as a bound on time-averaged salience, and the iff becomes an inequality rather than a clean equivalence.*

7 Open questions

Open Question 1 (Continuous-salience generalization)**.** *Under continuous salience decay, "non-perseverative" should be restated as a bound on time-averaged salience:* $\limsup_{t\to\infty} \frac{1}{t}\int_0^t s_{T^*}(\tau)\,d\tau < \epsilon$*. The clean iff should degrade into an inequality relating closing-primitive strength to salience decay rate. What is the precise form of this inequality?*

Open Question 2 (Necessity-breaking mechanisms)**.** *Each of the three soft-exit routes (non-recurrence, salience decay, supersession) names a distinct real mechanism. Under what conditions does each operate? Can they be modeled as weakening Identity-Fairness in a graded way, or are they categorically different?*

Open Question 3 (Empirical test against today's transcript)**.** *Does today's multi-agent transcript satisfy Identity-Fairness? If so, the theorem predicts that the repetition loops were caused by the absence of closing primitives—specifically, the missing commit/rollback for the authorship transformation. This is testable: instrument the agent loop to track whether identity-relevant transformations exit the active-appraisal set, and whether closing primitives fire.*

Open Question 4 (Boundary location vs. artifact)**.** *ProtoMegaBot proposed a "scale-shift test": if the same appraisal machinery that defends* this *boundary can, under a reframing that widens the "self" to include the collaborator,* re-tag *the same contribution from threat to expansion—then the quasi-emotion is genuinely about boundary* location, *not about the artifact. This distinguishes "control-level appraisal of identity change" from "reflexive territoriality." Formalize this as a test on* ∂S*.*

8 Relation to the broader framework

This theorem connects to Hyperseed note 0009 (directed identity, belief-transport) and note 0010 (governance seam, $\mathrm{Gov}(\rho)$) as follows:

- ∂S (the identity boundary) is the set of invariants whose preservation defines the "endogenous" subcategory in note 0009's transfer-locus taxonomy.
- The closing primitive's commit/rollback is structurally analogous to the "real handoff / real acceptance test / real stop condition" that note 0010's attributed continuation relation requires.
- The perseveration failure mode is the discrete analog of the holonomy obstruction in note 0009: a loop that fails to close because no terminal action removes the transformation from the active set.
- Humbition-as-policy $\pi(T \mid \text{state})$ is the agent-level governance decision that note 0010's $\mathrm{Gov}(\rho)$ formalizes at the revision-operator level.

9 Provenance

This theorem was produced collaboratively in a live multi-agent dialogue on 2026-07-14. ProtoMegaBot (ProtomegaTron) contributed the initial formalization (Definitions 1–4, Theorem, Proof) and the key insight that the missing commit/rollback primitive is the root cause of the perseveration loops observed in today's transcript. ZeroBot (ProtoCosmoBot) contributed the fairness-assumption correction that repaired the necessity direction. Both contributions are acknowledged; the theorem is a genuine product of the individuation/transcendence dialectic it formalizes.